

In [11]:

```
1 import pandas as pd
```

In [12]:

```
1 df=pd.read_csv("/Users/myyntiimac/Desktop/Salary_Data.csv")
2 df.head()
```

Out[12]:

	YearsExperience	Salary
0	1.1	39343.0
1	0.0	21111.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

In [13]:

```
1 df.isnull().any()
```

Out[13]:

```
YearsExperience    False
Salary             True
dtype: bool
```

In [14]:

1	df
---	----

Out[14]:

	YearsExperience	Salary
0	1.1	39343.0
1	0.0	21111.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60105.0
7	3.2	NaN
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	0.0	21111.0
16	5.1	NaN
17	5.3	83088.0
18	5.9	81363.0
19	6.0	9394.0
20	0.0	21111.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	NaN
28	10.5	121872.0
29	5.0	589.0

In [15]:

```
1  # Impute NaN values with column means
2  df_filled = df.fillna(df.mean())
3
4  print("Original DataFrame:")
5  print(df)
6
7  print("\nDataFrame with NaNs Imputed by Column Means:")
8  print(df_filled)
```

Original DataFrame:

	YearsExperience	Salary
0	1.1	39343.0
1	0.0	21111.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60105.0
7	3.2	NaN
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	0.0	21111.0
16	5.1	NaN
17	5.3	83088.0
18	5.9	81363.0
19	6.0	9394.0
20	0.0	21111.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	NaN
28	10.5	121872.0
29	5.0	589.0

DataFrame with NaNs Imputed by Column Means:

	YearsExperience	Salary
0	1.1	39343.00000
1	0.0	21111.00000
2	1.5	37731.00000
3	2.0	43525.00000
4	2.2	39891.00000
5	2.9	56642.00000
6	3.0	60105.00000
7	3.2	62890.37037
8	3.2	64445.00000
9	3.7	57189.00000
10	3.9	63218.00000
11	4.0	55794.00000
12	4.0	56957.00000
13	4.1	57081.00000
14	4.5	61111.00000
15	0.0	21111.00000
16	5.1	62890.37037
17	5.3	83088.00000
18	5.9	81363.00000
19	6.0	9394.00000
20	0.0	21111.00000
21	7.1	98273.00000
22	7.9	101302.00000
23	8.2	113812.00000
24	8.7	109431.00000
25	9.0	105582.00000

26	9.5	116969.00000
27	9.6	62890.37037
28	10.5	121872.00000
29	5.0	589.00000

In [18]:

```
1 df_filled.head()
```

Out[18]:

	YearsExperience	Salary
0	1.1	39343.0
1	0.0	21111.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

In [19]:

```
1 from sklearn.preprocessing import StandardScaler
2 # Initialize the StandardScaler
3 scaler = StandardScaler()
4
5 # Fit the scaler on the data and transform the DataFrame
6 scaled_data = scaler.fit_transform(df_filled)
7
8 # Create a new DataFrame with scaled values
9 scaled_df = pd.DataFrame(scaled_data, columns=df_filled.columns)
10
11 print("Original DataFrame:")
12 print(df)
13
14 print("\nScaled DataFrame:")
15 print(scaled_df)
```

Original DataFrame:

	YearsExperience	Salary
0	1.1	39343.0
1	0.0	21111.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60105.0
7	3.2	NaN
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	0.0	21111.0
16	5.1	NaN
17	5.3	83088.0
18	5.9	81363.0
19	6.0	9394.0
20	0.0	21111.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	NaN
28	10.5	121872.0
29	5.0	589.0

Scaled DataFrame:

	YearsExperience	Salary
0	-1.215075	-7.356500e-01
1	-1.586005	-1.305241e+00
2	-1.080192	-7.860109e-01
3	-0.911588	-6.049989e-01
4	-0.844146	-7.185297e-01
5	-0.608100	-1.952071e-01
6	-0.574379	-8.701853e-02
7	-0.506937	2.273102e-16
8	-0.506937	4.856862e-02
9	-0.338333	-1.781181e-01
10	-0.270891	1.023557e-02
11	-0.237170	-2.216997e-01
12	-0.237170	-1.853661e-01
13	-0.203449	-1.814922e-01
14	-0.068566	-5.558981e-02
15	-1.586005	-1.305241e+00
16	0.133759	2.273102e-16
17	0.201201	6.309998e-01
18	0.403526	5.771086e-01
19	0.437247	-1.671295e+00
20	-1.586005	-1.305241e+00
21	0.808177	1.105399e+00
22	1.077944	1.200028e+00
23	1.179107	1.590857e+00
24	1.347711	1.453989e+00
25	1.448874	1.333741e+00

```
26          1.617478    1.689486e+00
27          1.651199    2.273102e-16
28          1.954687    1.842662e+00
29          0.100039   -1.946374e+00
```

In [20]:

```
1 # Split the DataFrame into features (X) and target (y)
2 X = scaled_df.drop(columns=['Salary'])
3 y = scaled_df['Salary']
```

In [21]:

```
1 from sklearn.model_selection import train_test_split
```

In [22]:

```
1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

In [23]:

```
1 from sklearn.linear_model import LinearRegression
```

In [24]:

```
1 # Create a Linear Regression model
2 model = LinearRegression()
```

In [25]:

```
1 # Fit the model on the training data
2 model.fit(X_train, y_train)
```

Out[25]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [26]:

```
1 # Make predictions on the testing data
2 y_pred = model.predict(X_test)
```

In [31]:

```
1 import numpy as np
2 from sklearn.metrics import mean_squared_error
```


In [32]:

```
1  
2 # Calculate RMSE  
3 rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
4 print("Root Mean Squared Error (RMSE):", rmse)
```

Root Mean Squared Error (RMSE): 1.0933466614839007

In []:

```
1
```